



CISPA
HELMHOLTZ CENTER FOR
INFORMATION SECURITY



TECHNISCHE
UNIVERSITÄT
WIEN



European Research Council
Established by the European Commission

Symbolic Bug Finding for Asynchronous Hyperproperties

Workshop on Asynchronous Hyperproperties

A. Correnson, T. Nießen, B. Finkbeiner, G. Weissenbacher
July 6, 2024

Part I: Motivation

Motivating Example

Let's look at a simple voting protocol for two candidates A and B ...

```
countA  $\leftarrow$  0  
countB  $\leftarrow$  0  
loop  
  input  $vote \in \{A, B\}$   
  if  $vote = A$  then  
    countA  $\leftarrow$  countA + 1  
  if  $vote = B$  then  
    countB  $\leftarrow$  countB + 1
```

Motivating Example

... there could be a *tiny bug* in its implementation that plays in the favor of candidate B (whoops)

```
countA  $\leftarrow$  0  
countB  $\leftarrow$  0  
loop  
  input vote  $\in \{A, B\}$   
  if vote =  $A$  then  
    countA  $\leftarrow$  countA + 1  
  if vote =  $B$  then  
    countB  $\leftarrow$  countA + 1
```

A subtle bug

Observations

Observations

1. This bug is surprisingly difficult to catch with traditional bug-finding tools without giving a full formal specification of the voting protocol.

Observations

1. This bug is surprisingly difficult to catch with traditional bug-finding tools without giving a full formal specification of the voting protocol.
2. Even without a precise specification, it is clear that this code cannot be correct because it does not treat candidates A and B equally.

Observations

1. This bug is surprisingly difficult to catch with traditional bug-finding tools without giving a full formal specification of the voting protocol.
2. Even without a precise specification, it is clear that this code cannot be correct because it does not treat candidates A and B equally.

"candidates A and B should be treated equally" is an example of a **hyperproperty**
 \implies it requires comparing several executions of the voting protocol

Specifying Hyperproperties in HyperLTL

To formally specify hyperproperties, we can use the logic [HyperLTL](#), an extension of LTL with trace quantification.

Specifying Hyperproperties in HyperLTL

To formally specify hyperproperties, we can use the logic [HyperLTL](#), an extension of LTL with trace quantification.

Trace quantification

$$\psi ::= \forall \tau. \psi \mid \exists \tau. \psi \mid \varphi$$

where φ is a temporal relation between traces

Specifying Hyperproperties in HyperLTL

To formally specify hyperproperties, we can use the logic [HyperLTL](#), an extension of LTL with trace quantification.

Trace quantification

$$\psi ::= \forall \tau. \psi \mid \exists \tau. \psi \mid \varphi$$

where φ is a temporal relation between traces

Temporal relations

$$\varphi ::= P \mid \varphi \wedge \varphi \mid \Box \varphi \mid \varphi \cup \varphi \mid \dots$$

where P is a predicate over program variables labeled by a trace variable τ

Going back to our example of the voting protocol, one way to formalize the intuition "candidates A and B should be treated equally" is by requiring for the protocol to be **symmetric**.

Going back to our example of the voting protocol, one way to formalize the intuition "candidates A and B should be treated equally" is by requiring for the protocol to be **symmetric**.

$$\forall \tau_1. \forall \tau_2. \Box(\text{vote}_{\tau_1} = A \leftrightarrow \text{vote}_{\tau_2} = B) \rightarrow \\ \Box(\text{count}A_{\tau_1} = \text{count}B_{\tau_2} \wedge \text{count}B_{\tau_1} = \text{count}A_{\tau_2})$$

Going back to our example of the voting protocol, one way to formalize the intuition "candidates A and B should be treated equally" is by requiring for the protocol to be **symmetric**.

$$\forall \tau_1. \forall \tau_2. \Box(\text{vote}_{\tau_1} = A \leftrightarrow \text{vote}_{\tau_2} = B) \rightarrow \\ \Box(\text{count}A_{\tau_1} = \text{count}B_{\tau_2} \wedge \text{count}B_{\tau_1} = \text{count}A_{\tau_2})$$

Another, simpler, way is to require that candidates have **equal opportunities**

Going back to our example of the voting protocol, one way to formalize the intuition "candidates A and B should be treated equally" is by requiring for the protocol to be **symmetric**.

$$\forall \tau_1. \forall \tau_2. \Box (vote_{\tau_1} = A \leftrightarrow vote_{\tau_2} = B) \rightarrow \\ \Box (countA_{\tau_1} = countB_{\tau_2} \wedge countB_{\tau_1} = countA_{\tau_2})$$

Another, simpler, way is to require that candidates have **equal opportunities**

$$\forall \tau_1. \exists \tau_2. \Box (countA_{\tau_1} = countB_{\tau_2} \wedge countA_{\tau_2} = countB_{\tau_1})$$

Going back to our example of the voting protocol, one way to formalize the intuition "candidates A and B should be treated equally" is by requiring for the protocol to be **symmetric**.

$$\forall \tau_1. \forall \tau_2. \Box(\text{vote}_{\tau_1} = A \leftrightarrow \text{vote}_{\tau_2} = B) \rightarrow \\ \Box(\text{count}A_{\tau_1} = \text{count}B_{\tau_2} \wedge \text{count}B_{\tau_1} = \text{count}A_{\tau_2})$$

Another, simpler, way is to require that candidates have **equal opportunities**

$$\forall \tau_1. \exists \tau_2. \Box(\text{count}A_{\tau_1} = \text{count}B_{\tau_2} \wedge \text{count}A_{\tau_2} = \text{count}B_{\tau_1})$$

Importantly, the **alternation of universal and existential quantification** enables a concise specification that does not refer to the inputs of the protocol

Synchronization Problems (i)

Consider the following slightly modified version of the voting protocol:

```
countA  $\leftarrow$  0  
countB  $\leftarrow$  0  
loop  
  input vote  $\in \{A, B\}$   
  if vote = A then  
    countA  $\leftarrow$  countA - 1  
    countA  $\leftarrow$  countA + 2  
  if vote = B then  
    countB  $\leftarrow$  countB + 1
```

Synchronization Problems (i)

Consider the following slightly modified version of the voting protocol:

```
countA  $\leftarrow$  0
countB  $\leftarrow$  0
loop
  input vote  $\in \{A, B\}$ 
  if vote = A then
    countA  $\leftarrow$  countA - 1
    countA  $\leftarrow$  countA + 2
  if vote = B then
    countB  $\leftarrow$  countB + 1
```

- This version is functionally equivalent to the initial protocol
- However, it trivially violates the **equal opportunities** property

Synchronization Problems (ii)

→ In the HyperLTL specification of **equal opportunities**, the invariant

$$\text{count}A_{\tau_1} = \text{count}B_{\tau_2} \wedge \text{count}A_{\tau_2} = \text{count}B_{\tau_1}$$

is required to hold **at every computation step**

Synchronization Problems (ii)

→ In the HyperLTL specification of **equal opportunities**, the invariant

$$\text{count}A_{\tau_1} = \text{count}B_{\tau_2} \wedge \text{count}A_{\tau_2} = \text{count}B_{\tau_1}$$

is required to hold **at every computation step**

→ Since updating the score for A takes one more step than updating the score for B , the invariant will always be violated temporarily

Synchronization Problems (ii)

- In the HyperLTL specification of **equal opportunities**, the invariant

$$\text{count}A_{\tau_1} = \text{count}B_{\tau_2} \wedge \text{count}A_{\tau_2} = \text{count}B_{\tau_1}$$

is required to hold **at every computation step**

- Since updating the score for A takes one more step than updating the score for B , the invariant will always be violated temporarily
- To avoid this problem, we need an **asynchronous** logic enabling to specify how the different executions are **aligned** before being compared

Specifying Asynchronous Hyperproperties in OHyperLTL

OHyperLTL extends HyperLTL with support for asynchronous reasoning

Specifying Asynchronous Hyperproperties in OHyperLTL

OHyperLTL extends HyperLTL with support for asynchronous reasoning

$$\psi ::= \forall \tau : o. \psi \mid \exists \tau : o. \psi \mid \varphi$$


where o is a predicate specifying which states should be **observed**

Specifying Asynchronous Hyperproperties in OHyperLTL

OHyperLTL extends HyperLTL with support for asynchronous reasoning

$$\psi ::= \forall \tau : o.\psi \mid \exists \tau : o.\psi \mid \varphi$$

where o is a predicate specifying which states should be **observed**



Example

```
ℓ0: countA ← 0
ℓ1: countB ← 0
ℓ2: loop
ℓ3:   input vote ∈ {A, B}
ℓ4:   if vote = A then
ℓ5:     countA ← countA − 1
ℓ6:     countA ← countA + 2
ℓ7:   if vote = B then
ℓ8:     countB ← countB + 1
```

$$\forall \tau_1 : @ \ell_3, \exists \tau_2 : @ \ell_3, \Box (countA_{\tau_1} = countB_{\tau_2} \wedge countA_{\tau_2} = countB_{\tau_1})$$

Towards automated bug-hunting for OHyperLTL?

Summary of observations made so far

Summary of observations made so far

1. Some subtle bugs can only be detected by checking programs against hyperproperties

Summary of observations made so far

1. Some subtle **bugs** can only be detected by checking programs against hyperproperties
2. **Alternation of \forall and \exists trace quantifiers** is a convenient/necessary feature for concise specification of relevant hyperproperties

Summary of observations made so far

1. Some subtle **bugs** can only be detected by checking programs against hyperproperties
2. **Alternation of \forall and \exists trace quantifiers** is a convenient/necessary feature for concise specification of relevant hyperproperties
3. In the context of software, hyperproperties we wish to specify are **asynchronous**

Towards automated bug-hunting for OHyperLTL?

Summary of observations made so far

1. Some subtle **bugs** can only be detected by checking programs against hyperproperties
2. **Alternation of \forall and \exists trace quantifiers** is a convenient/necessary feature for concise specification of relevant hyperproperties
3. In the context of software, hyperproperties we wish to specify are **asynchronous**

Our goal

→ a **fully automated bug-hunting** technique for $\forall^*\exists^*$ asynchronous hyperproperties expressed in OHyperLTL

Verification is extremely difficult

- requires finding a proof that, for every first trace, there exists a second trace that satisfies the specified relation

Verification is extremely difficult

- requires finding a proof that, for every first trace, there exists a second trace that satisfies the specified relation

Refutation is not (much) simpler

- requires finding a trace and a proof that, for this trace, no second trace exists that satisfies the specified relation

Part II: Symbolic Execution for Asynchronous Hyperproperties

Many existing approaches even for $\forall\exists$ hyperproperties, but. . .

- Game-based verification
 - typically incomplete and generally does not produce counterexamples

Many existing approaches even for $\forall\exists$ hyperproperties, but. . .

- Game-based verification
→ typically incomplete and generally does not produce counterexamples
- Hoare-style relational verification
→ requires expert guidance (loop invariants, predicate abstractions, . . .)

Many existing approaches even for $\forall\exists$ hyperproperties, but. . .

- Game-based verification
→ typically incomplete and generally does not produce counterexamples
- Hoare-style relational verification
→ requires expert guidance (loop invariants, predicate abstractions, . . .)
- Automata-based model-checking and QBF-based bounded model-checking
→ limited to the analysis of finite-state systems

Many existing approaches even for $\forall\exists$ hyperproperties, but. . .

- Game-based verification
→ typically incomplete and generally does not produce counterexamples
- Hoare-style relational verification
→ requires expert guidance (loop invariants, predicate abstractions, . . .)
- Automata-based model-checking and QBF-based bounded model-checking
→ limited to the analysis of finite-state systems

Many existing approaches even for $\forall\exists$ hyperproperties, but. . .

- Game-based verification
→ typically incomplete and generally does not produce counterexamples
 - Hoare-style relational verification
→ requires expert guidance (loop invariants, predicate abstractions, . . .)
 - Automata-based model-checking and QBF-based bounded model-checking
→ limited to the analysis of finite-state systems
- no existing approach can fully automatically find counterexamples to $\forall\exists$ hyperproperties in asynchronous, infinite-state systems


Taking a step back: Symbolic execution for traditional bug-finding

Symbolic execution explores all feasible paths within a program by computing a symbolic encoding of the program's behavior.

Taking a step back: Symbolic execution for traditional bug-finding

Symbolic execution explores all feasible paths within a program by computing a symbolic encoding of the program's behavior.


```
while  $x > 0$   
   $x := x + 1$   
  assert  $x < 3$ 
```



Taking a step back: Symbolic execution for traditional bug-finding

Symbolic execution explores all feasible paths within a program by computing a symbolic encoding of the program's behavior.

- logical representation of program paths
- reduces bug finding to SMT solving
- symbolic encoding must be computable



Symbolic Execution for OHyperLTL_{safe}

Given: OHyperLTL_{safe} hyperproperty $\forall \tau_1. \exists \tau_2. \Box \varphi$

Idea: use two symbolic execution engines to find a model for $\exists \tau_1. \forall \tau_2. \neg \Box \varphi$

Symbolic Execution for OHyperLTL_{safe}

Given: OHyperLTL_{safe} hyperproperty $\forall \tau_1. \exists \tau_2. \Box \varphi$

Idea: use two symbolic execution engines to find a model for $\exists \tau_1. \forall \tau_2. \neg \Box \varphi$

- one symbolic execution engine searches for candidate symbolic paths π_1
(for a bounded number of observation points)

Symbolic Execution for OHyperLTL_{safe}

Given: OHyperLTL_{safe} hyperproperty $\forall \tau_1. \exists \tau_2. \Box \varphi$

Idea: use **two symbolic execution engines** to find a model for $\exists \tau_1. \forall \tau_2. \neg \Box \varphi$

- one symbolic execution engine searches for candidate symbolic paths π_1 (for a bounded number of observation points)
- a second symbolic execution engine tries to refute that π_1 is a counterexample by searching for a path π_2 such that φ holds for π_1, π_2

Symbolic Execution for OHyperLTL_{safe}

Given: OHyperLTL_{safe} hyperproperty $\forall \tau_1. \exists \tau_2. \Box \varphi$

Idea: use **two symbolic execution engines** to find a model for $\exists \tau_1. \forall \tau_2. \neg \Box \varphi$

- one symbolic execution engine searches for candidate symbolic paths π_1 (for a bounded number of observation points)
 - a second symbolic execution engine tries to refute that π_1 is a counterexample by searching for a path π_2 such that φ holds for π_1, π_2
- when this refutation fails, we have found a hyperbug!

Symbolic Execution for OHyperLTL_{safe}

Given: OHyperLTL_{safe} hyperproperty $\forall \tau_1. \exists \tau_2. \Box \varphi$

Idea: use **two symbolic execution engines** to find a model for $\exists \tau_1. \forall \tau_2. \neg \Box \varphi$

- one symbolic execution engine searches for candidate symbolic paths π_1 (for a bounded number of observation points)
 - a second symbolic execution engine tries to refute that π_1 is a counterexample by searching for a path π_2 such that φ holds for π_1, π_2
- when this refutation fails, we have found a hyperbug!
- otherwise, continue with next candidate π_1

Symbolic Execution for OHyperLTL_{safe}

Given: OHyperLTL_{safe} hyperproperty $\forall \tau_1. \exists \tau_2. \Box \varphi$

Idea: use **two symbolic execution engines** to find a model for $\exists \tau_1. \forall \tau_2. \neg \Box \varphi$

- one symbolic execution engine searches for candidate symbolic paths π_1 (for a bounded number of observation points)
 - a second symbolic execution engine tries to refute that π_1 is a counterexample by searching for a path π_2 such that φ holds for π_1, π_2
- when this refutation fails, we have found a hyperbug!
- otherwise, continue with next candidate π_1

Symbolic Execution for OHyperLTL_{safe}

Given: OHyperLTL_{safe} hyperproperty $\forall \tau_1. \exists \tau_2. \Box \varphi$

Idea: use **two symbolic execution engines** to find a model for $\exists \tau_1. \forall \tau_2. \neg \Box \varphi$

- one symbolic execution engine searches for candidate symbolic paths π_1 (for a bounded number of observation points)
 - a second symbolic execution engine tries to refute that π_1 is a counterexample by searching for a path π_2 such that φ holds for π_1, π_2
- when this refutation fails, we have found a hyperbug!
- otherwise, continue with next candidate π_1

Can be generalized to $\forall^* \exists^* \dots$ hyperproperties through product constructions

Problem solved?! No. . .

What about asynchronous properties in infinite-state systems?

Problem solved?! No...

What about asynchronous properties in infinite-state systems?

Soundness requires the second symbolic execution engine to consider all possible symbolic paths (for the fixed number of observation points).

Problem solved?! No...

What about asynchronous properties in infinite-state systems?

Soundness requires the second symbolic execution engine to consider all possible symbolic paths (for the fixed number of observation points).

This is not generally possible in asynchronous infinite-state systems, hence this approach is **incomplete**.

Problem solved?! No. . .

What about asynchronous properties in infinite-state systems?

Soundness requires the second symbolic execution engine to consider all possible symbolic paths (for the fixed number of observation points).

This is not generally possible in asynchronous infinite-state systems, hence this approach is **incomplete**.

For $\forall\exists$ hyperproperties, **computable symbolic encodings** do not generally exist!

Lemma

*Refuting an asynchronous $\forall\exists$ hyperproperty of an infinite-state system is **undecidable**.*

By reduction from the halting problem.

Lemma

*Refuting an asynchronous $\forall\exists$ hyperproperty of an infinite-state system is **undecidable**.*

By reduction from the halting problem.

Corollary

*Any **sound** bug finding algorithm for asynchronous $\forall\exists$ hyperproperties of infinite-state systems is necessarily **incomplete**.*

Relative completeness for observable programs

Definition (Informal)

A program is **observable** if, in any symbolic state, there is some n such that, after n steps, the program either produces an observed system state or terminates.

Relative completeness for observable programs

Definition (Informal)

A program is **observable** if, in any symbolic state, there is some n such that, after n steps, the program either produces an observed system state or terminates.

Theorem

$\text{OHyperLTL}_{\text{safe}}$ permits a **computable symbolic encoding** $\llbracket \psi \rrbracket_{\mathcal{T}}$ for **observable** input programs over some theory \mathcal{T} .

Relative completeness for observable programs

Definition (Informal)

A program is **observable** if, in any symbolic state, there is some n such that, after n steps, the program either produces an observed system state or terminates.

Theorem

$\text{OHyperLTL}_{\text{safe}}$ permits a **computable symbolic encoding** $\llbracket \psi \rrbracket_{\mathcal{T}}$ for **observable** input programs over some theory \mathcal{T} .

Corollary

Symbolic execution-based refutation of $\text{OHyperLTL}_{\text{safe}}$ hyperproperties is **complete** for **observable** input programs, assuming \mathcal{T} is decidable.

Part III: Evaluation

Experimental results - ORHLE benchmarks

Class	Type	Program	FO	Bug found	# Combinations	Runtime
$\forall\exists$	Other	draw-once	✓	✓	1	0.001 s
$\forall\exists$	Refinement	simple-nonrefinement	✓	✓	1	0.001 s
$\forall\exists$	Other	do-nothing	✓	✓	1	0.001 s
$\forall\forall\exists$	Generalized non-interference	nondet-leak2	✓	✓	2	0.001 s
$\forall\forall\exists$	Generalized non-interference	simple-leak	✓	✓	1	0.001 s
$\forall\forall\exists$	Generalized non-interference	smith1	✓	✓	2	0.003 s
$\forall\forall\exists$	Generalized non-interference	nondet-leak	✓	✓	2	0.003 s
$\forall\forall\exists$	Delimited release	parity-no-dr	✓	✓	2	0.003 s
$\forall\forall\exists$	Delimited release	wallet-no-dr	✓	✓	2	0.003 s
$\forall\exists$	Refinement	conditional-nonrefinement	✓	✓	4	0.006 s
$\forall\exists$	Refinement	add3-shuffled	✓	✓	6	0.009 s
$\forall\forall\forall\exists$	Delimited release	conditional-no-dr	✓	✓	8	0.013 s
$\forall\forall\exists$	Delimited release	median-no-dr	✓	✓	4	0.016 s
$\forall\forall\forall\exists$	Generalized non-interference	conditional-leak	✓	✓	48	0.074 s
$\forall\exists$	Refinement	loop-nonrefinement	✗	✗	N/A	∞

Experimental results - our own benchmarks

Class	Program	Bug found	# Observations	# Combinations	Runtime
$\forall\exists$	even_odd	✓	1	1	0.001 s
$\forall\exists$	factor2	✓	2	2	0.001 s
$\forall\exists$	for_loop_simple	✓	1	2	0.010 s
$\forall\exists$	linear_equation	✓	22	22	0.023 s
$\forall\exists$	monotonic_increase	✓	7	7	0.029 s
$\forall\exists$	escalating	✓	7	747	0.103 s
$\forall\forall\exists$	secret_pin_leak	✓	8	11	0.103 s
$\forall\exists$	escalating_2	✓	7	1707	0.190 s
$\forall\forall$	obs_determinism	✓	4	86	0.408 s
$\forall\exists$	no_primes_above_31397	✓	1	201	1.203 s
$\forall\forall\exists$	secret_pin_leak_2	✓	3	248	1.972 s
$\forall\exists$	exponential_branching_1	✓	1	1024	2.216 s
$\forall\exists$	exponential_branching_2	✓	1	2048	4.040 s