



**CISPA**

HELMHOLTZ CENTER FOR  
INFORMATION SECURITY



TECHNISCHE  
UNIVERSITÄT  
WIEN



European Research Council  
Established by the European Commission

# Automated Hyperbug Finding

A. Correnson, T. Nießen, B. Finkbeiner, G. Weissenbacher  
July 23, 2024

# Part I: Motivation

## Motivating Example

Let's look at a simple voting protocol for two candidates  $A$  and  $B$ ...

```
 $\ell_0$ :  $countA \leftarrow 0$   
 $\ell_1$ :  $countB \leftarrow 0$   
 $\ell_2$ : loop  
 $\ell_3$ :   input  $vote \in \{A, B\}$   
 $\ell_4$ :   if  $vote = A$  then  
 $\ell_5$ :      $countA \leftarrow countA + 1$   
 $\ell_6$ :   if  $vote = B$  then  
 $\ell_7$ :      $countB \leftarrow countB + 1$ 
```

## Motivating Example

... there could be a *tiny bug* in its implementation that plays in the favor of candidate  $B$  (whoops)

```
l0: countA ← 0
l1: countB ← 0
l2: loop
l3:   input vote ∈ {A, B}
l4:   if vote = A then
l5:     countA ← countA + 1
l6:   if vote = B then
l7:     countB ← countA + 1
```

# A subtle bug

## Observations

## Observations

1. This bug is surprisingly difficult to catch with traditional bug-finding tools without giving a full formal specification of the voting protocol.

## Observations

1. This bug is surprisingly difficult to catch with traditional bug-finding tools without giving a full formal specification of the voting protocol.
2. Even without a precise specification, it is clear that this code cannot be correct because it does not treat candidates  $A$  and  $B$  equally.

## Observations

1. This bug is surprisingly difficult to catch with traditional bug-finding tools without giving a full formal specification of the voting protocol.
2. Even without a precise specification, it is clear that this code cannot be correct because it does not treat candidates  $A$  and  $B$  equally.

"candidates  $A$  and  $B$  should be treated equally" is an example of a **hyperproperty**  
 $\implies$  it requires comparing several executions of the voting protocol



## Specifying Hyperproperties in HyperLTL

To formally specify hyperproperties of software systems, we use the logic [HyperLTL](#), an extension of LTL with trace quantification.

# Specifying Hyperproperties in HyperLTL

To formally specify hyperproperties of software systems, we use the logic [HyperLTL](#), an extension of LTL with trace quantification.

## Trace quantification

$$\psi ::= \forall \tau. \psi \mid \exists \tau. \psi \mid \varphi$$

where  $\varphi$  is a temporal relation between traces

# Specifying Hyperproperties in HyperLTL

To formally specify hyperproperties of software systems, we use the logic [HyperLTL](#), an extension of LTL with trace quantification.

## Trace quantification

$$\psi ::= \forall \tau. \psi \mid \exists \tau. \psi \mid \varphi$$

where  $\varphi$  is a temporal relation between traces

## Temporal relations

$$\varphi ::= P \mid \varphi \wedge \varphi \mid \Box \varphi \mid \varphi \cup \varphi \mid \dots$$

where  $P$  is a predicate over program variables labeled by a trace variable  $\tau$

# Specifying Asynchronous Hyperproperties in OHyperLTL

OHyperLTL extends HyperLTL with support for asynchronous reasoning by introducing explicit observation points

# Specifying Asynchronous Hyperproperties in OHyperLTL

OHyperLTL extends HyperLTL with support for asynchronous reasoning by introducing explicit observation points

$$\psi ::= \forall \tau : o. \psi \mid \exists \tau : o. \psi \mid \varphi$$


where  $o$  is a predicate specifying which states should be **observed**

# Specifying Asynchronous Hyperproperties in OHyperLTL

OHyperLTL extends HyperLTL with support for asynchronous reasoning by introducing explicit observation points

$$\psi ::= \forall \tau : o.\psi \mid \exists \tau : o.\psi \mid \varphi$$

where  $o$  is a predicate specifying which states should be **observed**



Going back to our example of the voting protocol, one way to formalize the intuition "candidates  $A$  and  $B$  should be treated equally" is by requiring for the protocol to be **symmetric**.

Going back to our example of the voting protocol, one way to formalize the intuition "candidates  $A$  and  $B$  should be treated equally" is by requiring for the protocol to be **symmetric**.

$$\forall \tau_1 : \ell_3. \forall \tau_2 : \ell_3. \Box(\text{vote}_{\tau_1} = A \leftrightarrow \text{vote}_{\tau_2} = B) \rightarrow \\ \Box(\text{count}A_{\tau_1} = \text{count}B_{\tau_2} \wedge \text{count}B_{\tau_1} = \text{count}A_{\tau_2})$$



Going back to our example of the voting protocol, one way to formalize the intuition "candidates  $A$  and  $B$  should be treated equally" is by requiring for the protocol to be **symmetric**.

$$\forall \tau_1 : \ell_3. \forall \tau_2 : \ell_3. \Box(\text{vote}_{\tau_1} = A \leftrightarrow \text{vote}_{\tau_2} = B) \rightarrow \\ \Box(\text{count}A_{\tau_1} = \text{count}B_{\tau_2} \wedge \text{count}B_{\tau_1} = \text{count}A_{\tau_2})$$

Another, simpler, way is to require that candidates have **equal opportunities**

Going back to our example of the voting protocol, one way to formalize the intuition "candidates  $A$  and  $B$  should be treated equally" is by requiring for the protocol to be **symmetric**.

$$\forall \tau_1 : \ell_3. \forall \tau_2 : \ell_3. \Box(\text{vote}_{\tau_1} = A \leftrightarrow \text{vote}_{\tau_2} = B) \rightarrow \\ \Box(\text{count}A_{\tau_1} = \text{count}B_{\tau_2} \wedge \text{count}B_{\tau_1} = \text{count}A_{\tau_2})$$

Another, simpler, way is to require that candidates have **equal opportunities**

$$\forall \tau_1 : \ell_3. \exists \tau_2 : \ell_3. \Box(\text{count}A_{\tau_1} = \text{count}B_{\tau_2} \wedge \text{count}A_{\tau_2} = \text{count}B_{\tau_1})$$

Going back to our example of the voting protocol, one way to formalize the intuition "candidates  $A$  and  $B$  should be treated equally" is by requiring for the protocol to be **symmetric**.

$$\forall \tau_1 : \ell_3. \forall \tau_2 : \ell_3. \Box (vote_{\tau_1} = A \leftrightarrow vote_{\tau_2} = B) \rightarrow \\ \Box (countA_{\tau_1} = countB_{\tau_2} \wedge countB_{\tau_1} = countA_{\tau_2})$$

Another, simpler, way is to require that candidates have **equal opportunities**

$$\forall \tau_1 : \ell_3. \exists \tau_2 : \ell_3. \Box (countA_{\tau_1} = countB_{\tau_2} \wedge countA_{\tau_2} = countB_{\tau_1})$$

Importantly, the **alternation of universal and existential quantification** enables a concise specification that does not refer to the inputs of the protocol

# Towards automated bug-hunting for OHyperLTL?

**Summary of observations made so far**

## Summary of observations made so far

1. Some subtle bugs can only be detected by checking programs against hyperproperties

## Summary of observations made so far

1. Some subtle **bugs** can only be detected by checking programs against hyperproperties
2. **Alternation of  $\forall$  and  $\exists$  trace quantifiers** is a convenient/necessary feature for concise specification of relevant hyperproperties

## Summary of observations made so far

1. Some subtle **bugs** can only be detected by checking programs against hyperproperties
2. **Alternation of  $\forall$  and  $\exists$  trace quantifiers** is a convenient/necessary feature for concise specification of relevant hyperproperties
3. In the context of software, hyperproperties we wish to specify are **asynchronous**

# Towards automated bug-hunting for OHyperLTL?

## Summary of observations made so far

1. Some subtle **bugs** can only be detected by checking programs against hyperproperties
2. **Alternation of  $\forall$  and  $\exists$  trace quantifiers** is a convenient/necessary feature for concise specification of relevant hyperproperties
3. In the context of software, hyperproperties we wish to specify are **asynchronous**

## Our goal

→ a **fully automated bug-hunting** technique for  $\forall^*\exists^*$  asynchronous hyperproperties expressed in OHyperLTL



## **Verification is extremely difficult**

- requires finding a proof that, for every first trace, there exists a second trace that satisfies the specified relation

## **Verification is extremely difficult**

- requires finding a proof that, for every first trace, there exists a second trace that satisfies the specified relation

## **Refutation is not (much) simpler**

- requires finding a trace and a proof that, for this trace, no second trace exists that satisfies the specified relation

Many existing approaches even for  $\forall\exists$  hyperproperties, but. . .

- Game-based verification  
→ incomplete and does not produce counterexamples

Many existing approaches even for  $\forall\exists$  hyperproperties, but. . .

- Game-based verification  
→ incomplete and does not produce counterexamples
- Hoare-style relational verification  
→ requires expert guidance (loop invariants, predicate abstractions, . . .)

Many existing approaches even for  $\forall\exists$  hyperproperties, but. . .

- Game-based verification  
→ incomplete and does not produce counterexamples
- Hoare-style relational verification  
→ requires expert guidance (loop invariants, predicate abstractions, . . . )
- Automata-based model-checking and QBF-based bounded model-checking  
→ limited to the analysis of finite-state systems

Many existing approaches even for  $\forall\exists$  hyperproperties, but. . .

- Game-based verification  
→ incomplete and does not produce counterexamples
- Hoare-style relational verification  
→ requires expert guidance (loop invariants, predicate abstractions, . . . )
- Automata-based model-checking and QBF-based bounded model-checking  
→ limited to the analysis of finite-state systems

Many existing approaches even for  $\forall\exists$  hyperproperties, but. . .

- Game-based verification  
→ incomplete and does not produce counterexamples
- Hoare-style relational verification  
→ requires expert guidance (loop invariants, predicate abstractions, . . . )
- Automata-based model-checking and QBF-based bounded model-checking  
→ limited to the analysis of finite-state systems

→ no existing approach can fully automatically find counterexamples to  $\forall\exists$  hyperproperties in asynchronous, infinite-state systems

## Part II: Symbolic Execution for Asynchronous Hyperproperties



## Taking a step back: Symbolic execution for traditional bug-finding

Symbolic execution explores all behavior of a program by computing a symbolic encoding of the program's paths.

## Taking a step back: Symbolic execution for traditional bug-finding


Symbolic execution explores all behavior of a program by computing a symbolic encoding of the program's paths.

```
while  $x > 0$   
   $x := x + 1$   
  assert  $x < 3$ 
```

# Taking a step back: Symbolic execution for traditional bug-finding

Symbolic execution explores all behavior of a program by computing a symbolic encoding of the program's paths.


```
while  $x > 0$   
   $x := x + 1$   
  assert  $x < 3$ 
```



# Taking a step back: Symbolic execution for traditional bug-finding

Symbolic execution explores all behavior of a program by computing a symbolic encoding of the program's paths.


- logical representation of program paths
- reduces bug finding to SMT solving



## Taking a step back: Symbolic execution for traditional bug-finding

Symbolic execution explores all behavior of a program by computing a symbolic encoding of the program's paths.

$0 < x \wedge x + 1 \geq 3$  is sat  
 $\Rightarrow$  bug found!



# Symbolic Execution for OHyperLTL<sub>safe</sub>

**Given:** OHyperLTL<sub>safe</sub> hyperproperty  $\forall \tau_1. \exists \tau_2. \Box \varphi$

**Idea:** use **two symbolic execution engines** to find a model for  $\exists \tau_1. \forall \tau_2. \neg \Box \varphi$

# Symbolic Execution for OHyperLTL<sub>safe</sub>

**Given:** OHyperLTL<sub>safe</sub> hyperproperty  $\forall \tau_1. \exists \tau_2. \Box \varphi$

**Idea:** use **two symbolic execution engines** to find a model for  $\exists \tau_1. \forall \tau_2. \neg \Box \varphi$

- one symbolic execution engine searches for candidate symbolic paths  $\pi_1$   
(for a bounded number of observation points)

**Given:** OHyperLTL<sub>safe</sub> hyperproperty  $\forall \tau_1. \exists \tau_2. \Box \varphi$

**Idea:** use **two symbolic execution engines** to find a model for  $\exists \tau_1. \forall \tau_2. \neg \Box \varphi$

- one symbolic execution engine searches for candidate symbolic paths  $\pi_1$  (for a bounded number of observation points)
- a second symbolic execution engine tries to refute that  $\pi_1$  is a counterexample by searching for a path  $\pi_2$  such that  $\Box \varphi$  holds for  $\pi_1, \pi_2$



**Given:** OHyperLTL<sub>safe</sub> hyperproperty  $\forall \tau_1. \exists \tau_2. \Box \varphi$

**Idea:** use **two symbolic execution engines** to find a model for  $\exists \tau_1. \forall \tau_2. \neg \Box \varphi$

- one symbolic execution engine searches for candidate symbolic paths  $\pi_1$  (for a bounded number of observation points)
  - a second symbolic execution engine tries to refute that  $\pi_1$  is a counterexample by searching for a path  $\pi_2$  such that  $\Box \varphi$  holds for  $\pi_1, \pi_2$
- when this refutation fails, we have found a hyperbug!

**Given:** OHyperLTL<sub>safe</sub> hyperproperty  $\forall \tau_1. \exists \tau_2. \Box \varphi$

**Idea:** use **two symbolic execution engines** to find a model for  $\exists \tau_1. \forall \tau_2. \neg \Box \varphi$

- one symbolic execution engine searches for candidate symbolic paths  $\pi_1$  (for a bounded number of observation points)
  - a second symbolic execution engine tries to refute that  $\pi_1$  is a counterexample by searching for a path  $\pi_2$  such that  $\Box \varphi$  holds for  $\pi_1, \pi_2$
- when this refutation fails, we have found a hyperbug!

**Given:** OHyperLTL<sub>safe</sub> hyperproperty  $\forall \tau_1. \exists \tau_2. \Box \varphi$

**Idea:** use **two symbolic execution engines** to find a model for  $\exists \tau_1. \forall \tau_2. \neg \Box \varphi$

- one symbolic execution engine searches for candidate symbolic paths  $\pi_1$  (for a bounded number of observation points)
  - a second symbolic execution engine tries to refute that  $\pi_1$  is a counterexample by searching for a path  $\pi_2$  such that  $\Box \varphi$  holds for  $\pi_1, \pi_2$
- when this refutation fails, we have found a hyperbug!

Can be generalized to  $\forall^* \exists^*$  hyperproperties through product constructions

## The good news

### Theorem (Soundness)

*Symbolic Execution for  $\text{OHyperLTL}_{\text{safe}}$  is a sound hyperbug finding method.*

## The good news

### Theorem (Soundness)

*Symbolic Execution for  $\text{OHyperLTL}_{\text{safe}}$  is a sound hyperbug finding method.*

## The bad news

### Lemma (Undecidability)

*Refuting asynchronous  $\forall\exists$  hyperproperties of infinite-state systems is **undecidable**.  
(reduction from the halting problem)*

## The good news

### Theorem (Soundness)

*Symbolic Execution for  $\text{OHyperLTL}_{\text{safe}}$  is a sound hyperbug finding method.*

## The bad news

### Lemma (Undecidability)

*Refuting asynchronous  $\forall\exists$  hyperproperties of infinite-state systems is **undecidable**.  
(reduction from the halting problem)*

### Corollary (Incompleteness)

*Symbolic Execution for  $\text{OHyperLTL}_{\text{safe}}$  is necessarily incomplete.*

# What is the problem?

**Intuition:** using symbolic execution to find a model for  $\exists \tau_1. \forall \tau_2. \neg \Box \varphi$  requires to enumerate all possible symbolic paths of a given observation length

# What is the problem?

**Intuition:** using symbolic execution to find a model for  $\exists \tau_1. \forall \tau_2. \neg \Box \varphi$  requires to enumerate all possible symbolic paths of a given observation length

The set of symbolic paths of **observational length**  $k$  is not necessarily finite!

$\implies$  we cannot always enumerate symbolic paths of observation length  $k$



# What is the problem?

**Intuition:** using symbolic execution to find a model for  $\exists \tau_1. \forall \tau_2. \neg \Box \varphi$  requires to enumerate all possible symbolic paths of a given observation length

The set of symbolic paths of **observational length**  $k$  is not necessarily finite!

$\Rightarrow$  we cannot always enumerate symbolic paths of observation length  $k$

```
ℓ0: loop  
ℓ1:   while(...) {...}  
ℓ2:   observe
```

Symbolic paths of observational lengths  $k$  are of the form  $(\ell_0(\ell_1)^*\ell_2)^k$  (i.e., there are infinitely many of them)

### Definition (Observable Programs)

A program is **observable** if, in any symbolic state, there is some  $n$  such that, after  $n$  steps, the program either produces an observed system state or terminates.

# Relative completeness for observable programs

## Definition (Observable Programs)

A program is **observable** if, in any symbolic state, there is some  $n$  such that, after  $n$  steps, the program either produces an observed system state or terminates.

## Theorem

*Relative Completeness Symbolic Execution for  $\text{OHyperLTL}_{\text{safe}}$  is a complete hyperbug finding method for **observable** programs (assuming symbolic paths can be expressed in a decidable first-order theory).*

## Part III: Evaluation

# Experimental results - ORHLE benchmarks

Class	Type	Program	FO	Bug found	# Combinations	Runtime
$\forall\exists$	Other	draw-once	✓	✓	1	0.001 s
$\forall\exists$	Refinement	simple-nonrefinement	✓	✓	1	0.001 s
$\forall\exists$	Other	do-nothing	✓	✓	1	0.001 s
$\forall\forall\exists$	Generalized non-interference	nondet-leak2	✓	✓	2	0.001 s
$\forall\forall\exists$	Generalized non-interference	simple-leak	✓	✓	1	0.001 s
$\forall\forall\exists$	Generalized non-interference	smith1	✓	✓	2	0.003 s
$\forall\forall\exists$	Generalized non-interference	nondet-leak	✓	✓	2	0.003 s
$\forall\forall\exists$	Delimited release	parity-no-dr	✓	✓	2	0.003 s
$\forall\forall\exists$	Delimited release	wallet-no-dr	✓	✓	2	0.003 s
$\forall\exists$	Refinement	conditional-nonrefinement	✓	✓	4	0.006 s
$\forall\exists$	Refinement	add3-shuffled	✓	✓	6	0.009 s
$\forall\forall\forall\exists$	Delimited release	conditional-no-dr	✓	✓	8	0.013 s
$\forall\forall\exists$	Delimited release	median-no-dr	✓	✓	4	0.016 s
$\forall\forall\forall\exists$	Generalized non-interference	conditional-leak	✓	✓	48	0.074 s
$\forall\exists$	Refinement	loop-nonrefinement	✗	✗	N/A	$\infty$

# Experimental results - our own benchmarks

Class	Program	Bug found	# Observations	# Combinations	Runtime
$\forall\exists$	even_odd	✓	1	1	0.001 s
$\forall\exists$	factor2	✓	2	2	0.001 s
$\forall\exists$	for_loop_simple	✓	1	2	0.010 s
$\forall\exists$	linear_equation	✓	22	22	0.023 s
$\forall\exists$	monotonic_increase	✓	7	7	0.029 s
$\forall\exists$	escalating	✓	7	747	0.103 s
$\forall\forall\exists$	secret_pin_leak	✓	8	11	0.103 s
$\forall\exists$	escalating_2	✓	7	1707	0.190 s
$\forall\forall$	obs_determinism	✓	4	86	0.408 s
$\forall\exists$	no_primes_above_31397	✓	1	201	1.203 s
$\forall\forall\exists$	secret_pin_leak_2	✓	3	248	1.972 s
$\forall\exists$	exponential_branching_1	✓	1	1024	2.216 s
$\forall\exists$	exponential_branching_2	✓	1	2048	4.040 s